

# Lab 6 - PCC177/BCC406

## REDES NEURAIS E APRENDIZAGEM EM PROFUNDIDADE

### Modelos Generativos

Prof. Eduardo e Prof. Pedro

Objetivos:

- Parte I : Compressão com AE
- Parte II : Detecção de anomalias
- Parte III: Redes Generativas Adversariais

Data da entrega : XX/XX

- Complete o código (marcado com `ToDo`) e quando requisitado, escreva textos diretamente nos notebooks. Onde tiver `None`, substitua pelo seu código.
- Execute todo notebook e salve tudo em um PDF **nomeado** como "NomeSobrenome-Lab.pdf"
- Envie o PDF via google [FORM](#)

Este notebook é baseado em tensorflow e Keras.

### Parte I: Autoencoder para redução de dimensionalidade (30pt)

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import tensorflow as tf

from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.model_selection import train_test_split
from tensorflow.keras import layers, losses
from tensorflow.keras.datasets import fashion_mnist
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.  
[Mostrar diferenças](#)

Carrega dataset Fashion MNIST dataset. Cada imagem tem resolução 28x28 pixels.

```
(x_train, _), (x_test, _) = fashion_mnist.load_data()

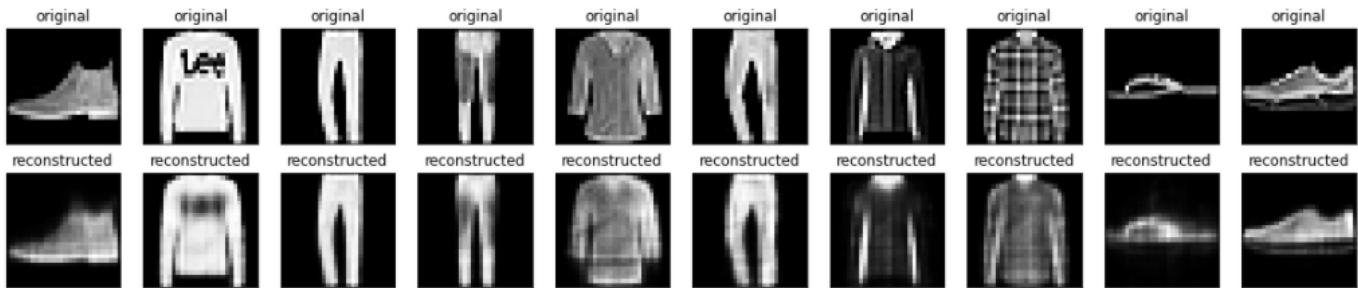
x_train = x_train.astype('float32') / 255.
```

```
x_test = x_test.astype('float32') / 255.

print (x_train.shape)
print (x_test.shape)

(60000, 28, 28)
(10000, 28, 28)
```

## ▼ Exemplo de classes



Abaixo exemplo de implementação de autoencoder apena com camadas densas. O encoder, comprime as imegns em 4 dimensões (latent\_dim), e o decoder reconstróe a imagem a partir do vetor latente.

O exemplo abaixo usa a [Keras Model Subclassing API](#).

```
latent_dim = 4

class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
        ])
        self.decoder = tf.keras.Sequential([
            layers.Dense(784, activation='sigmoid'),
            layers.Reshape((28, 28))
    ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)  
autoencoder = Autoencoder(latent\_dim)

```
autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())
```

```
autoencoder.fit(x_train, x_train,
```

```
epochs=10,  
shuffle=True,  
validation_data=(x_test, x_test))  
  
Epoch 1/10  
1875/1875 [=====] - 7s 2ms/step - loss: 0.0504 - val_loss:  
Epoch 2/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0359 - val_loss:  
Epoch 3/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0337 - val_loss:  
Epoch 4/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0331 - val_loss:  
Epoch 5/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0328 - val_loss:  
Epoch 6/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0327 - val_loss:  
Epoch 7/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0326 - val_loss:  
Epoch 8/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0326 - val_loss:  
Epoch 9/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0325 - val_loss:  
Epoch 10/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0325 - val_loss:  
<keras.callbacks.History at 0x7f86a01c4990>
```

Treine o modelo e veja os resultados da re-construção.

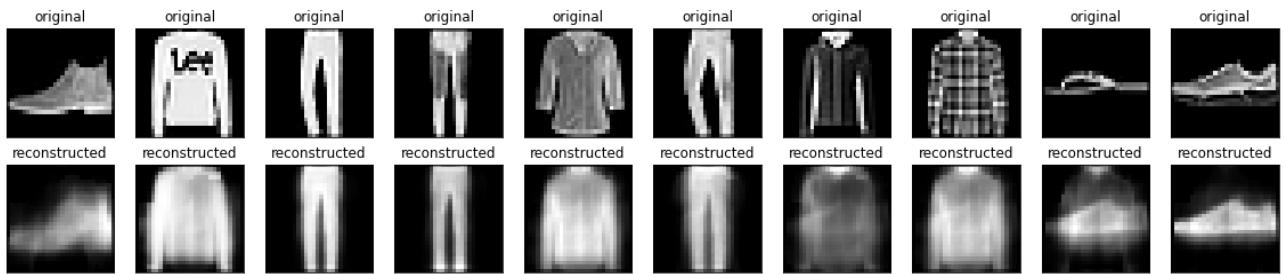
```
encoded_imgs = autoencoder.encoder(x_test).numpy()  
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()
```

```
n = 10  
plt.figure(figsize=(20, 4))  
for i in range(n):  
    # display original  
    ax = plt.subplot(2, n, i + 1)  
    plt.imshow(x_test[i])  
    plt.title("original")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)  
  
    # display reconstruction  
    ax = plt.subplot(2, n, i + 1 + n)  
    plt.imshow(decoded_imgs[i])
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
ax.get_xaxis().set_visible(False)  
ax.get_yaxis().set_visible(False)  
plt.show()
```



## ▼ ToDo : Testes (15pt)

Faça testes com vetor latente de dimensões 2, 8, 16 e 64.

Dim=2

```
latent_dim = 2

class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
        ])
        self.decoder = tf.keras.Sequential([
            layers.Dense(784, activation='sigmoid'),
            layers.Reshape((28, 28))
        ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded

autoencoder = Autoencoder(latent_dim)

autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
epochs=10,
shuffle=True,
validation_data=(x_test, x_test))
```

Epoch 1/10

1875/1875 [=====] - 4s 2ms/step - loss: 0.0688 - val\_loss:

Epoch 2/10

```
1875/1875 [=====] - 4s 2ms/step - loss: 0.0548 - val_loss:  
Epoch 3/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0486 - val_loss:  
Epoch 4/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0454 - val_loss:  
Epoch 5/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0446 - val_loss:  
Epoch 6/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0444 - val_loss:  
Epoch 7/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0443 - val_loss:  
Epoch 8/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0442 - val_loss:  
Epoch 9/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0442 - val_loss:  
Epoch 10/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0442 - val_loss:  
<keras.callbacks.History at 0x7f863612ac50>
```



```
encoded_imgs = autoencoder.encoder(x_test).numpy()  
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()
```

```
n = 10  
plt.figure(figsize=(20, 4))  
for i in range(n):  
    # display original  
    ax = plt.subplot(2, n, i + 1)  
    plt.imshow(x_test[i])  
    plt.title("original")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)  
  
    # display reconstruction  
    ax = plt.subplot(2, n, i + 1 + n)  
    plt.imshow(decoded_imgs[i])  
    plt.title("reconstructed")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)  
plt.show()
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)



Dim=8

```
reconstructed    reconstructed    reconstructed    reconstructed    reconstructed    reconstructed    reconstructed    reconstructed    reconstructed    reconstructed
```

latent\_dim = 8

```
class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
        ])
        self.decoder = tf.keras.Sequential([
            layers.Dense(784, activation='sigmoid'),
            layers.Reshape((28, 28))
    ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded
```

autoencoder = Autoencoder(latent\_dim)

autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())

```
autoencoder.fit(x_train, x_train,
                 epochs=10,
                 shuffle=True,
                 validation_data=(x_test, x_test))
```

```
Epoch 1/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0407 - val_loss:
Epoch 2/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0251 - val_loss:
Epoch 3/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0237 - val_loss:
Epoch 4/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0234 - val_loss:
Epoch 5/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0232 - val_loss:
Epoch 6/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0231 - val_loss:
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
Epoch 8/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0230 - val_loss:
Epoch 9/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0230 - val_loss:
Epoch 10/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0230 - val_loss:
<keras.callbacks.History at 0x7f86204ff350>
```

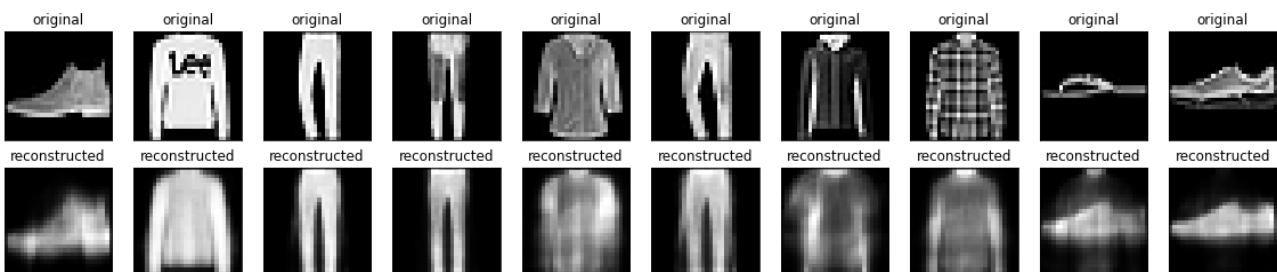
```

encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i])
    plt.title("original")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i])
    plt.title("reconstructed")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()

```



Dim=16

```

latent_dim = 16

class Autoencoder(Model):

```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```

    self.latent_dim = latent_dim
    self.encoder = tf.keras.Sequential([
        layers.Flatten(),
        layers.Dense(latent_dim, activation='relu'),
    ])
    self.decoder = tf.keras.Sequential([

```

```
        layers.Dense(784, activation='sigmoid'),
        layers.Reshape((28, 28))
    ])

def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

autoencoder = Autoencoder(latent_dim)

autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())

autoencoder.fit(x_train, x_train,
                 epochs=10,
                 shuffle=True,
                 validation_data=(x_test, x_test))

Epoch 1/10
1875/1875 [=====] - 5s 2ms/step - loss: 0.0346 - val_loss:
Epoch 2/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0199 - val_loss:
Epoch 3/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0182 - val_loss:
Epoch 4/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0179 - val_loss:
Epoch 5/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0177 - val_loss:
Epoch 6/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0176 - val_loss:
Epoch 7/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0175 - val_loss:
Epoch 8/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0175 - val_loss:
Epoch 9/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0174 - val_loss:
Epoch 10/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.0174 - val_loss:
<keras.callbacks.History at 0x7f86203ce190>
```

```
encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()
```

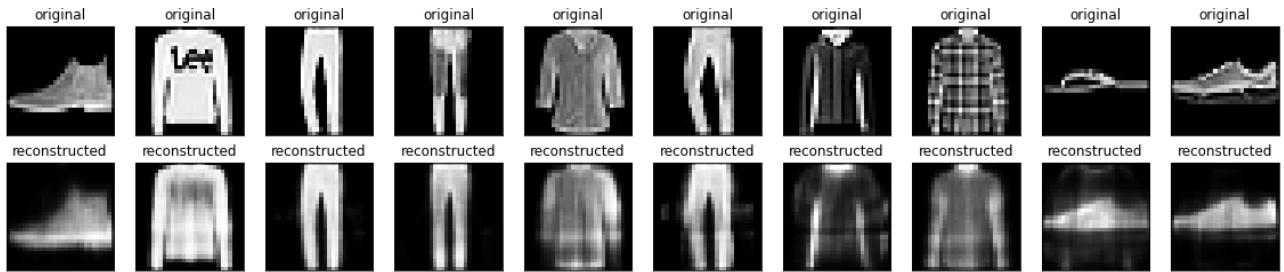
```
n = 10
plt.figure(figsize=(20, 4))
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
ax = plt.subplot(2, n, i + 1)
plt.imshow(x_test[i])
plt.title("original")
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
```

```
# display reconstruction
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i])
plt.title("reconstructed")
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```



Dim=64

```
latent_dim = 64

class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
        ])
        self.decoder = tf.keras.Sequential([
            layers.Dense(784, activation='sigmoid'),
            layers.Reshape((28, 28))
        ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())
```

```
autoencoder.fit(x_train, x_train,
                 epochs=10,
```

```
shuffle=True,  
validation_data=(x_test, x_test))  
  
Epoch 1/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0237 - val_loss:  
Epoch 2/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0116 - val_loss:  
Epoch 3/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0102 - val_loss:  
Epoch 4/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0096 - val_loss:  
Epoch 5/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0092 - val_loss:  
Epoch 6/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0091 - val_loss:  
Epoch 7/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0090 - val_loss:  
Epoch 8/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0089 - val_loss:  
Epoch 9/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0088 - val_loss:  
Epoch 10/10  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0088 - val_loss:  
<keras.callbacks.History at 0x7f8636304d10>
```

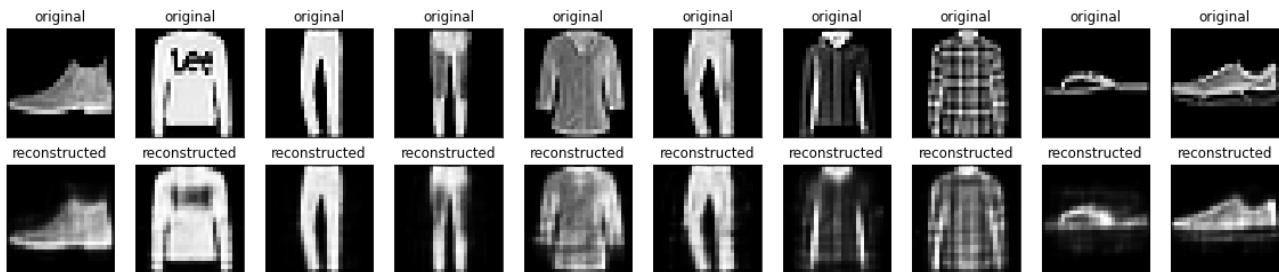


```
encoded_imgs = autoencoder.encoder(x_test).numpy()  
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()
```

```
n = 10  
plt.figure(figsize=(20, 4))  
for i in range(n):  
    # display original  
    ax = plt.subplot(2, n, i + 1)  
    plt.imshow(x_test[i])  
    plt.title("original")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)  
  
    # display reconstruction  
    ax = plt.subplot(2, n, i + 1 + n)  
    plt.imshow(decoded_imgs[i])  
    plt.title("reconstructed")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)



## ▼ ToDo : Responda (15pt)

Escreva suas conclusões sobre os testes executados:

Com base nos testes executados, podemos concluir que com o aumento do vetor latente ocorre uma melhora na resolução da imagem e tambem uma interferencia menor entre as imagens geradas.

## ▼ Parte II: Detecção de anomalias (30pt)

### Intro

Neste exemplo, você vai detectar anomalias em sinais de eletrocardiograma (ECG). Para tal, treine um autoencoder no dataset [ECG5000 dataset](#). Este dataset contém 5000 batimentos de ECG (<https://en.wikipedia.org/wiki/Electrocardiography>), cada um com 140 amostras (pontos) na curva. Cada instância da base de dados (um batimento) foi rotulado como zero (0) ou um (1). A classe zero corresponde a um batimento anormal e a classe um a um batimento de classe normal. Queremos identificar os anormais.

Para detectar anomalias usando um autoencoder você deve treinar um autoencoder apenas em batimentos normais. Ele vai aprender a re-construir os batimentos saudáveis. A hipótese é que os batimentos anormais vão divergir no padrão, quando compararmos a entrada com a reconstrução.

## ▼ Carrega base de ECG

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
# Download the dataset
dataframe = pd.read_csv('http://storage.googleapis.com/download.tensorflow.org/data/ecg.csv')
raw_data = dataframe.values
dataframe.head()
```

	0	1	2	3	4	5	6	7
0	-0.112522	-2.827204	-3.773897	-4.349751	-4.376041	-3.474986	-2.181408	-1.818286
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022377	-3.234368	-1.566126	-0.992258
2	-0.567088	-2.593450	-3.874230	-4.584095	-4.187449	-3.151462	-1.742940	-1.490659
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268016	-3.881110	-2.993280	-1.671131
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338224	-3.802422	-2.534510	-1.783423

5 rows × 141 columns



```
# The last element contains the labels
labels = raw_data[:, -1]

# The other data points are the electrocadriogram data
data = raw_data[:, 0:-1]

train_data, test_data, train_labels, test_labels = train_test_split(
    data, labels, test_size=0.2, random_state=21
)
```

Normaliza entre [0,1].

```
min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)
```

Vamos separar os batimentos normais (label 1) para treinar o Autoencoder.

```
train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

normal_train_data = train_data[train_labels]
normal_test_data = test_data[test_labels]
```

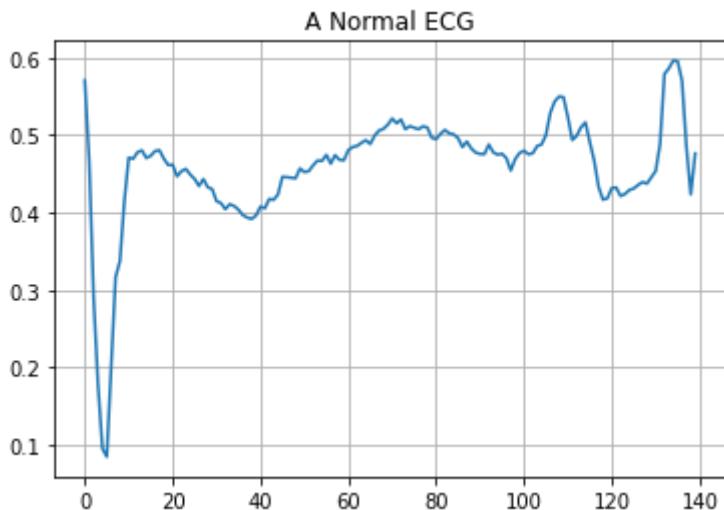
Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
anomalous_test_data = test_data[~test_labels]
```

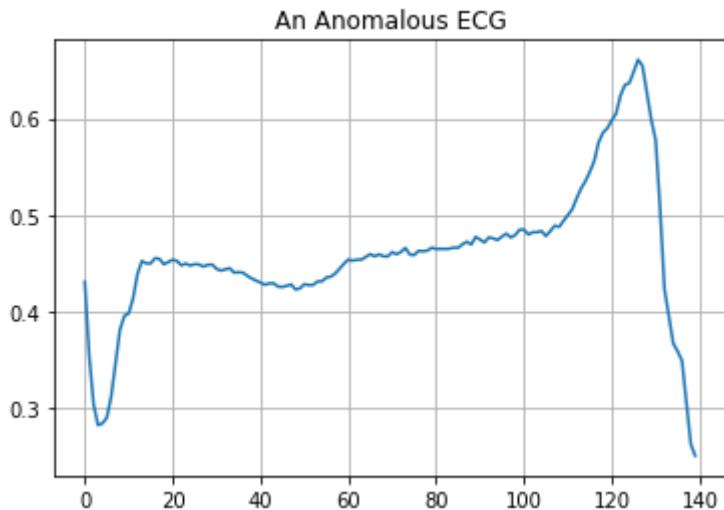
Plote um batimento normal.

```
plt.grid()
plt.plot(np.arange(140), normal_train_data[0])
plt.title("A Normal ECG")
plt.show()
```



Plote um batimento anômalo.

```
plt.grid()
plt.plot(np.arange(140), anomalous_train_data[0])
plt.title("An Anomalous ECG")
plt.show()
```



## ▼ ToDo : Construção de um modelo (30pt)

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

Tente construir um modelo com camadas de convolução de uma dimensão (Lembre-se que um sinal de ECG é uma série temporal de uma dimensão). [Conv1D](#)

```
latent_dim = 8
```

```

class AnomalyDetector(Model):
    def __init__(self):
        super(AnomalyDetector, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
        ])
        self.decoder = tf.keras.Sequential([
            layers.Dense(140, activation='sigmoid'),
            #layers.Reshape((140,0))
        ])
    #self.encoder = tf.keras.Sequential([
    #    # Todo: crie o encoder, com um gargalo

    #self.decoder = tf.keras.Sequential([
    #    # Todo: crie as camadas do decoder

def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

autoencoder = AnomalyDetector()

```

```
autoencoder.compile(optimizer='adam', loss='mae')
```

Depois de treinar com os batimentos normais, avalie com os anormais.

```

history = autoencoder.fit(normal_train_data, normal_train_data,
                           epochs=20,
                           batch_size=512,
                           validation_data=(test_data, test_data),
                           shuffle=True)

Epoch 1/20
5/5 [=====] - 1s 29ms/step - loss: 0.0571 - val_loss: 0.052
Epoch 2/20
5/5 [=====] - 0s 9ms/step - loss: 0.0547 - val_loss: 0.0510
Epoch 3/20
5/5 [=====] - 0s 12ms/step - loss: 0.0524 - val_loss: 0.049
Epoch 4/20
5/5 [=====] - 0s 10ms/step - loss: 0.0500 - val_loss: 0.048
Epoch 5/20
5/5 [=====] - 0s 8ms/step - loss: 0.0474 - val_loss: 0.0469

```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

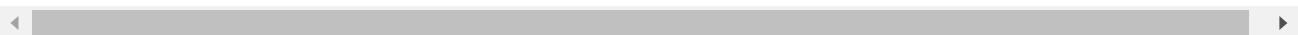
[Mostrar diferenças](#)

```

Epoch 6/20
5/5 [=====] - 0s 9ms/step - loss: 0.0420 - val_loss: 0.0442
Epoch 8/20
5/5 [=====] - 0s 8ms/step - loss: 0.0394 - val_loss: 0.0428
Epoch 9/20
5/5 [=====] - 0s 8ms/step - loss: 0.0369 - val_loss: 0.0415
Epoch 10/20

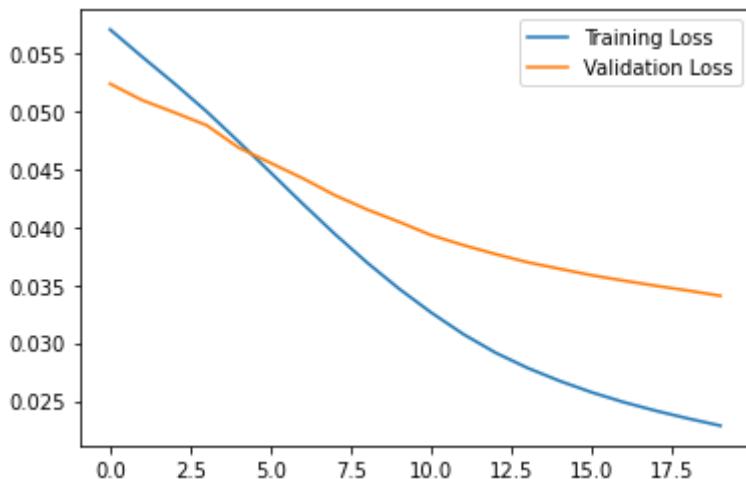
```

```
5/5 [=====] - 0s 8ms/step - loss: 0.0347 - val_loss: 0.0405
Epoch 11/20
5/5 [=====] - 0s 9ms/step - loss: 0.0326 - val_loss: 0.0393
Epoch 12/20
5/5 [=====] - 0s 12ms/step - loss: 0.0308 - val_loss: 0.038
Epoch 13/20
5/5 [=====] - 0s 8ms/step - loss: 0.0292 - val_loss: 0.0377
Epoch 14/20
5/5 [=====] - 0s 8ms/step - loss: 0.0278 - val_loss: 0.0370
Epoch 15/20
5/5 [=====] - 0s 12ms/step - loss: 0.0267 - val_loss: 0.036
Epoch 16/20
5/5 [=====] - 0s 8ms/step - loss: 0.0257 - val_loss: 0.0358
Epoch 17/20
5/5 [=====] - 0s 8ms/step - loss: 0.0249 - val_loss: 0.0354
Epoch 18/20
5/5 [=====] - 0s 8ms/step - loss: 0.0241 - val_loss: 0.0349
Epoch 19/20
5/5 [=====] - 0s 10ms/step - loss: 0.0235 - val_loss: 0.034
Epoch 20/20
5/5 [=====] - 0s 8ms/step - loss: 0.0229 - val_loss: 0.0341
```



```
plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()
```

<matplotlib.legend.Legend at 0x7f85fcdda990>

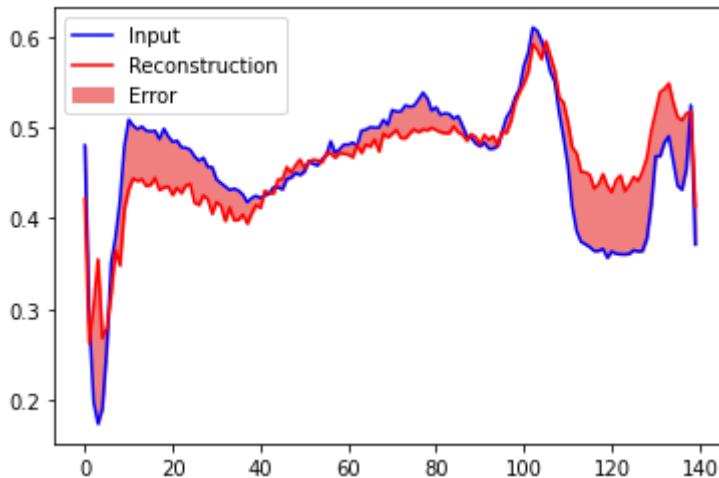


Você vai considerar um batimento como anômalo se ele divergir mais que um desvio padrão das amostras normais. Primeiro, vamos plotar um batimento normal a partir da base de treino e sua reconstrução. Assim, poderemos calcular o erro de re-construção.

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

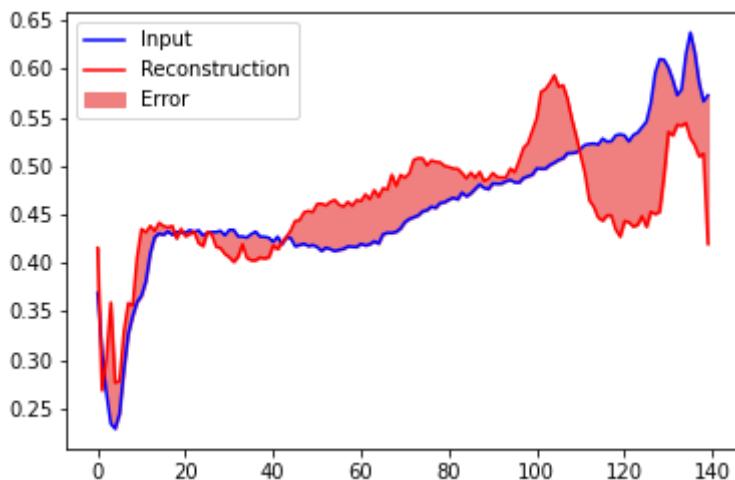
```
plt.plot(normal_test_data[0], 'b')
plt.plot(decoded_data[0], 'r')
plt.fill_between(np.arange(140), decoded_data[0], normal_test_data[0], color='lightcoral')
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



Vamos fazer o mesmo para um batimento anômalo.

```
encoded_data = autoencoder.encoder(anomalous_test_data).numpy()
decoded_data = autoencoder.decoder(encoded_data).numpy()

plt.plot(anomalous_test_data[0], 'b')
plt.plot(decoded_data[0], 'r')
plt.fill_between(np.arange(140), decoded_data[0], anomalous_test_data[0], color='lightcoral')
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



## ▼ Detectando as anomalias

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

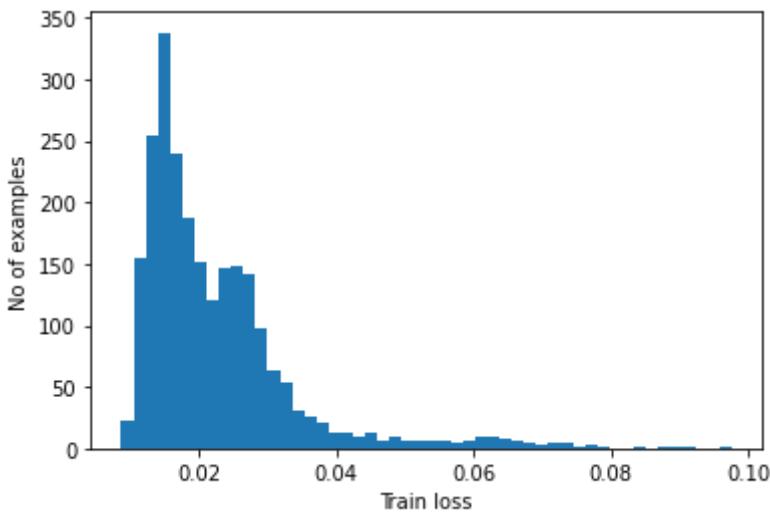
[Mostrar diferenças](#)

calcular o erro médio para os exemplos normais do treino e depois, classificar os anormais do teste, que tenha erro de reconstrução maior que um desvio padrão.

Plota erro de reconstrução de batimentos normais do treino

```
reconstructions = autoencoder.predict(normal_train_data)
train_loss = tf.keras.losses.mae(reconstructions, normal_train_data)

plt.hist(train_loss[None,:], bins=50)
plt.xlabel("Train loss")
plt.ylabel("No of examples")
plt.show()
```



Escolha do limiar.

```
threshold = np.mean(train_loss) + np.std(train_loss)
print("Threshold: ", threshold)

Threshold:  0.034112945

reconstructions = autoencoder.predict(anomalous_test_data)
test_loss = tf.keras.losses.mae(reconstructions, anomalous_test_data)

plt.hist(test_loss[None, :], bins=50)
plt.xlabel("Test loss")
plt.ylabel("No of examples")
plt.show()
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

Classificação.

```
def predict(model, data, threshold):
    reconstructions = model(data)
    loss = tf.keras.losses.mae(reconstructions, data)
    return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
    print("Accuracy = {}".format(accuracy_score(labels, predictions)))
    print("Precision = {}".format(precision_score(labels, predictions)))
    print("Recall = {}".format(recall_score(labels, predictions)))
```

Calcule a acurácia para os dois modelos (com camadas densas e convolucionais)

```
preds = predict(autoencoder, test_data, threshold)
print_stats(preds, test_labels)
```

```
Accuracy = 0.943
Precision = 0.9941060903732809
Recall = 0.9035714285714286
```

## ▼ Parte III: Redes Generativas Adversariais (40pt)

Leia o tutorial sobre a pix2pix em [Tensorflow Tutorials](#). O pix2pix foi apresentado em [Image-to-image translation with conditional adversarial networks by Isola et al. \(2017\)](#) e se trata de uma rede generativa adversarial condicional para geração de fachadas de prédios condicionada a uma máscara representando a arquitetura. baixe o notebook do tutorial, estude e treine a GAN. Após o treinamento, construa você mesmo 3 máscaras (usando algum software de desenho) e faça uma inferência com a rede. Anexe no notebook a máscara e sua respectiva saída.

## ▼ ToDo : Fachadas de prédios (40pt)

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
import pathlib
import time
import datetime

from matplotlib import pyplot as plt
from IPython import display
```

```
dataset_name = "facades"

_URL = f'http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/{dataset_name}.tar.gz'

path_to_zip = tf.keras.utils.get_file(
    fname=f'{dataset_name}.tar.gz',
    origin=_URL,
    extract=True)

path_to_zip = pathlib.Path(path_to_zip)

PATH = path_to_zip.parent/dataset_name

Downloading data from http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
30171136/30168306 [=====] - 19s 1us/step
30179328/30168306 [=====] - 19s 1us/step

◀ ▶
```

```
list(PATH.parent.iterdir())

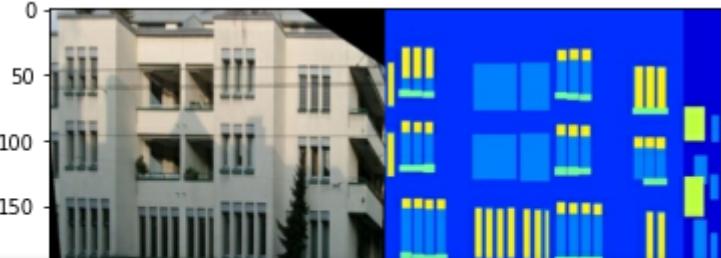
[PosixPath('/root/.keras/datasets/facades'),
 PosixPath('/root/.keras/datasets/fashion-mnist'),
 PosixPath('/root/.keras/datasets/facades.tar.gz')]

sample_image = tf.io.read_file(str(PATH / 'train/1.jpg'))
sample_image = tf.io.decode_jpeg(sample_image)
print(sample_image.shape)

(256, 512, 3)

plt.figure()
plt.imshow(sample_image)

<matplotlib.image.AxesImage at 0x7f85fc1cfa10>
```



Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.  
[Mostrar diferenças](#)

```
def load(image_file):
    # Read and decode an image file to a uint8 tensor
    image = tf.io.read_file(image_file)
```

```
image = tf.io.decode_jpeg(image)

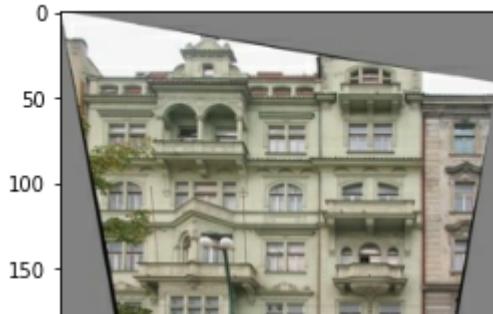
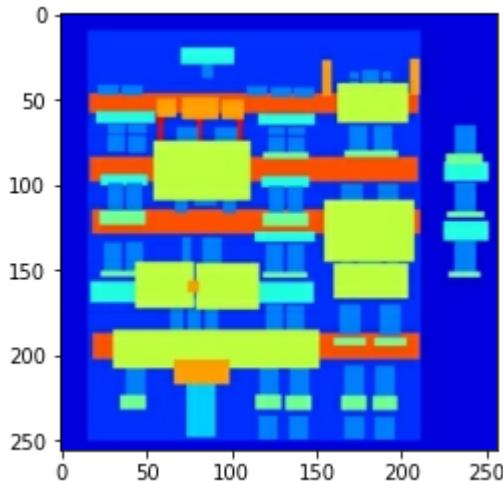
# Split each image tensor into two tensors:
# - one with a real building facade image
# - one with an architecture label image
w = tf.shape(image)[1]
w = w // 2
input_image = image[:, w:, :]
real_image = image[:, :w, :]

# Convert both images to float32 tensors
input_image = tf.cast(input_image, tf.float32)
real_image = tf.cast(real_image, tf.float32)

return input_image, real_image
```

```
inp, re = load(str(PATH / 'train/100.jpg'))
# Casting to int for matplotlib to display the images
plt.figure()
plt.imshow(inp / 255.0)
plt.figure()
plt.imshow(re / 255.0)
```

```
<matplotlib.image.AxesImage at 0x7f85fcc70710>
```



Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)



```
# The facade training set consist of 400 images
BUFFER_SIZE = 400
```

```
# The batch size of 1 produced better results for the U-Net in the original pix2pix experi
BATCH_SIZE = 1
# Each image is 256x256 in size
IMG_WIDTH = 256
IMG_HEIGHT = 256

def resize(input_image, real_image, height, width):
    input_image = tf.image.resize(input_image, [height, width],
                                  method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    real_image = tf.image.resize(real_image, [height, width],
                                 method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)

    return input_image, real_image

def random_crop(input_image, real_image):
    stacked_image = tf.stack([input_image, real_image], axis=0)
    cropped_image = tf.image.random_crop(
        stacked_image, size=[2, IMG_HEIGHT, IMG_WIDTH, 3])

    return cropped_image[0], cropped_image[1]

# Normalizing the images to [-1, 1]
def normalize(input_image, real_image):
    input_image = (input_image / 127.5) - 1
    real_image = (real_image / 127.5) - 1

    return input_image, real_image

@tf.function()
def random_jitter(input_image, real_image):
    # Resizing to 286x286
    input_image, real_image = resize(input_image, real_image, 286, 286)

    # Random cropping back to 256x256
    input_image, real_image = random_crop(input_image, real_image)

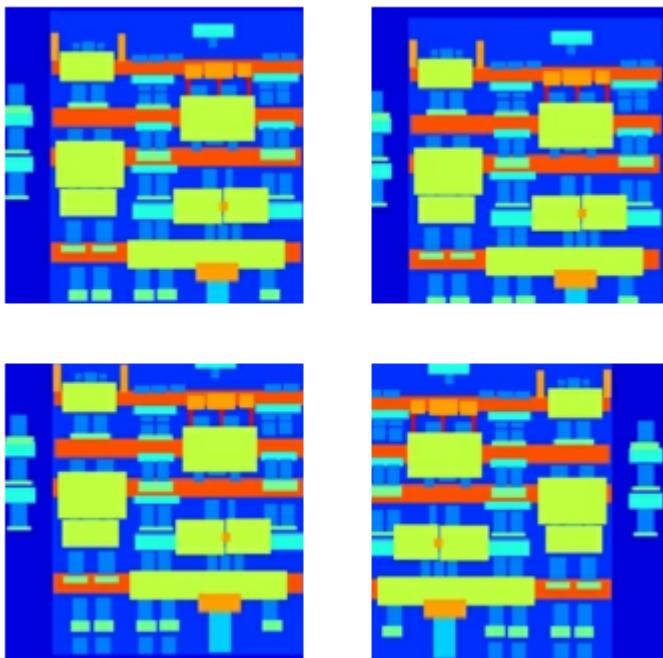
    if tf.random.uniform(() > 0.5:
        # Random mirroring
        input_image = tf.image.flip_left_right(input_image)
        real_image = tf.image.flip_left_right(real_image)

    return input_image, real_image
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
rj_inp = np.array(rj_in, np.float32)
plt.subplot(2, 2, i + 1)
plt.imshow(rj_inp / 255.0)
plt.axis('off')
plt.show()
```



```

def load_image_train(image_file):
    input_image, real_image = load(image_file)
    input_image, real_image = random_jitter(input_image, real_image)
    input_image, real_image = normalize(input_image, real_image)

    return input_image, real_image

def load_image_test(image_file):
    input_image, real_image = load(image_file)
    input_image, real_image = resize(input_image, real_image,
                                    IMG_HEIGHT, IMG_WIDTH)
    input_image, real_image = normalize(input_image, real_image)

    return input_image, real_image

train_dataset = tf.data.Dataset.list_files(str(PATH / 'train/*.jpg'))
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)

try:
    test_dataset = tf.data.Dataset.list_files(str(PATH / 'test/*.jpg'))
except tf.errors.InvalidArgumentError:
    test_dataset = tf.data.Dataset.list_files(str(PATH / 'val/*.jpg'))
test_dataset = test_dataset.map(load_image_test)

```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

OUTPUT\_CHANNELS = 3

```

def downsample(filters, size, apply_batchnorm=True):
    initializer = tf.random_normal_initializer(0., 0.02)

```

```
result = tf.keras.Sequential()
result.add(
    tf.keras.layers.Conv2D(filters, size, strides=2, padding='same',
                          kernel_initializer=initializer, use_bias=False))

if apply_batchnorm:
    result.add(tf.keras.layers.BatchNormalization())

result.add(tf.keras.layers.LeakyReLU())

return result

down_model = downsample(3, 4)
down_result = down_model(tf.expand_dims(inp, 0))
print (down_result.shape)

(1, 128, 128, 3)

def upsample(filters, size, apply_dropout=False):
    initializer = tf.random_normal_initializer(0., 0.02)

    result = tf.keras.Sequential()
    result.add(
        tf.keras.layers.Conv2DTranspose(filters, size, strides=2,
                                       padding='same',
                                       kernel_initializer=initializer,
                                       use_bias=False))

    result.add(tf.keras.layers.BatchNormalization())

    if apply_dropout:
        result.add(tf.keras.layers.Dropout(0.5))

    result.add(tf.keras.layers.ReLU())

    return result

up_model = upsample(3, 4)
up_result = up_model(down_result)
print (up_result.shape)

(1, 256, 256, 3)
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
down_stack = [
    downsample(64, 4, apply_batchnorm=False), # (batch_size, 128, 128, 64)
    downsample(128, 4), # (batch_size, 64, 64, 128)
    downsample(256, 4), # (batch_size, 32, 32, 256)
    downsample(512, 4), # (batch_size, 16, 16, 512)
```

```
downsample(512, 4), # (batch_size, 8, 8, 512)
downsample(512, 4), # (batch_size, 4, 4, 512)
downsample(512, 4), # (batch_size, 2, 2, 512)
downsample(512, 4), # (batch_size, 1, 1, 512)
]

up_stack = [
    upsample(512, 4, apply_dropout=True), # (batch_size, 2, 2, 1024)
    upsample(512, 4, apply_dropout=True), # (batch_size, 4, 4, 1024)
    upsample(512, 4, apply_dropout=True), # (batch_size, 8, 8, 1024)
    upsample(512, 4), # (batch_size, 16, 16, 1024)
    upsample(256, 4), # (batch_size, 32, 32, 512)
    upsample(128, 4), # (batch_size, 64, 64, 256)
    upsample(64, 4), # (batch_size, 128, 128, 128)
]

initializer = tf.random_normal_initializer(0., 0.02)
last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS, 4,
                                      strides=2,
                                      padding='same',
                                      kernel_initializer=initializer,
                                      activation='tanh') # (batch_size, 256, 256, 3)

x = inputs

# Downsampling through the model
skips = []
for down in down_stack:
    x = down(x)
    skips.append(x)

skips = reversed(skips[:-1])

# Upsampling and establishing the skip connections
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = tf.keras.layers.concatenate([x, skip])

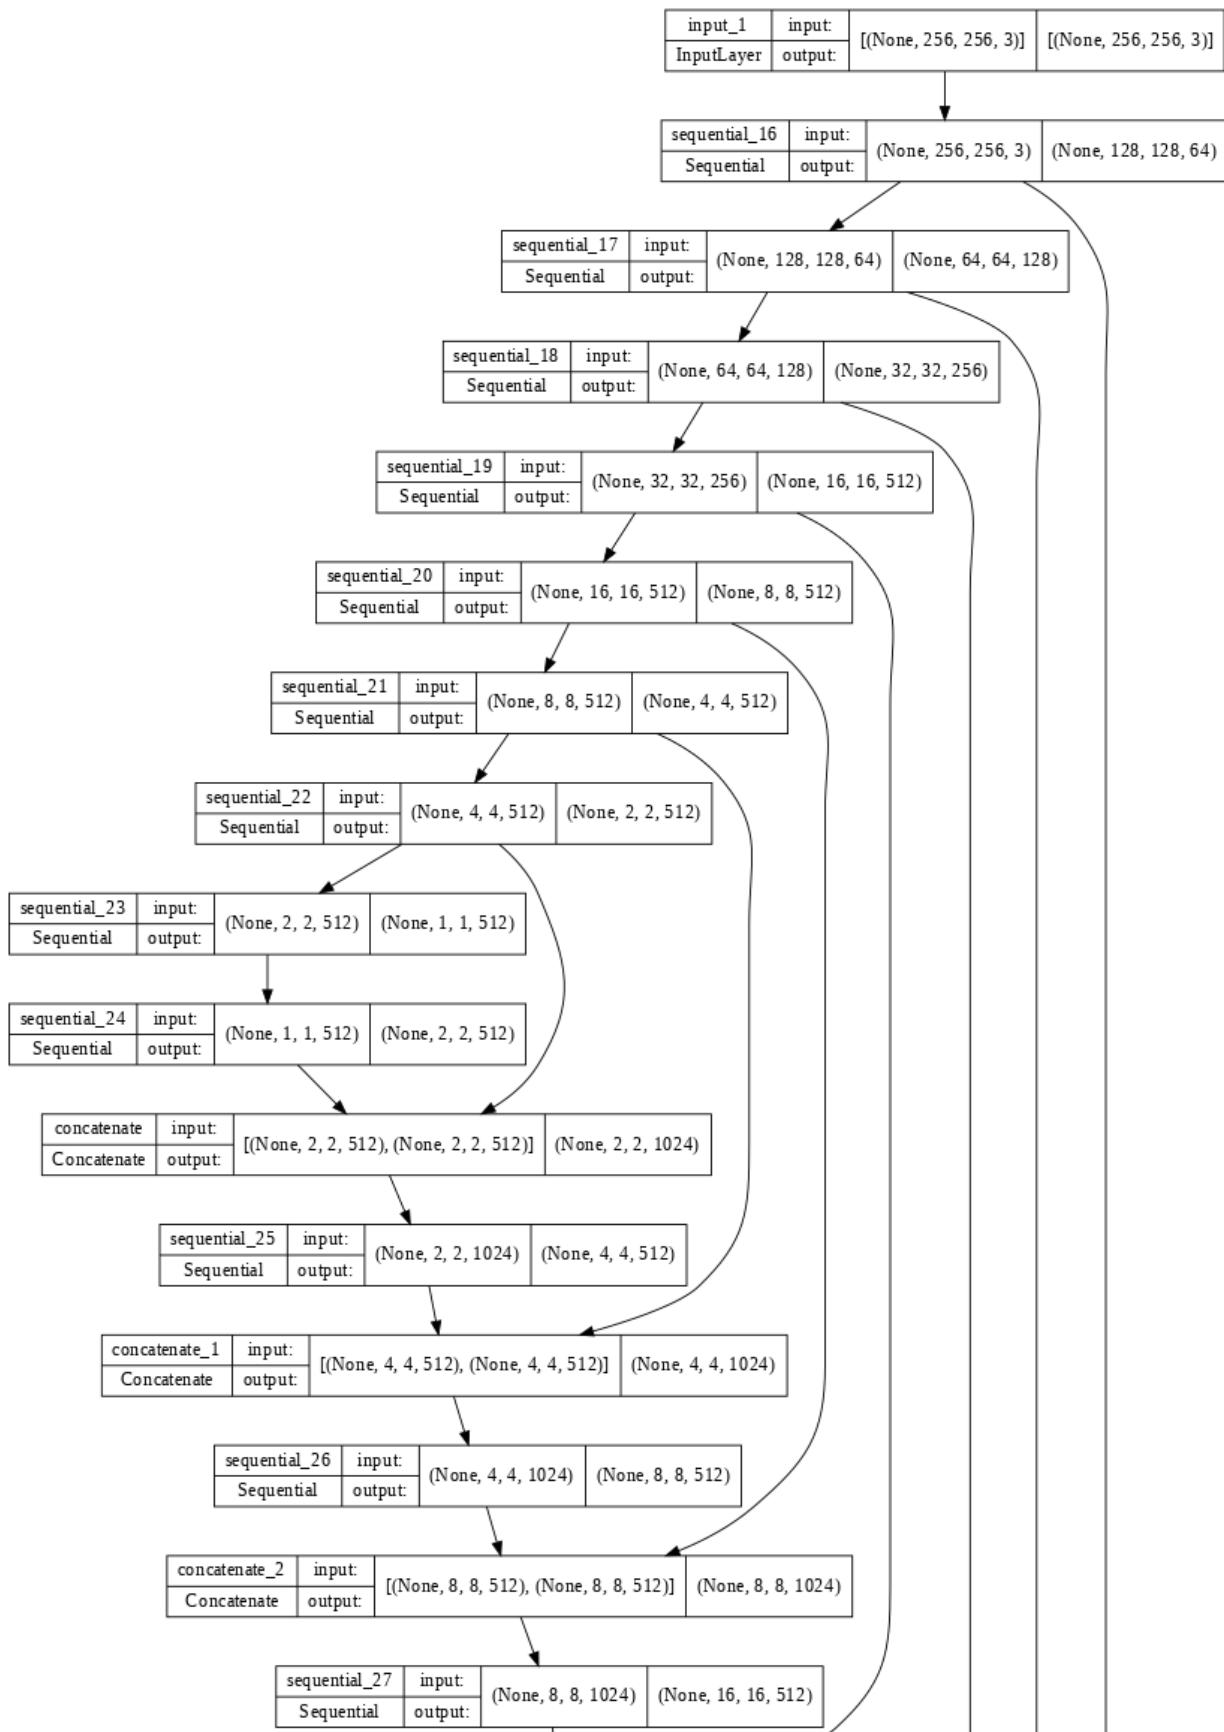
x = last(x)

return tf.keras.Model(inputs=inputs, outputs=x)

generator = Generator()
tf.keras.utils.plot_model(generator, show_shapes=True, dpi=64)
```

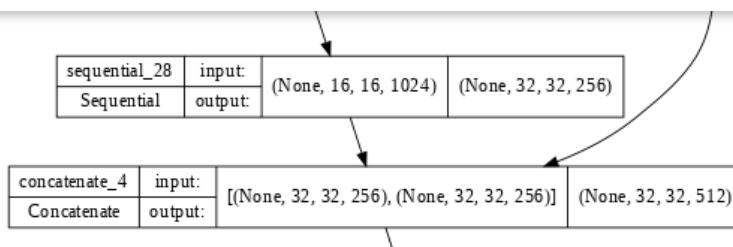
Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

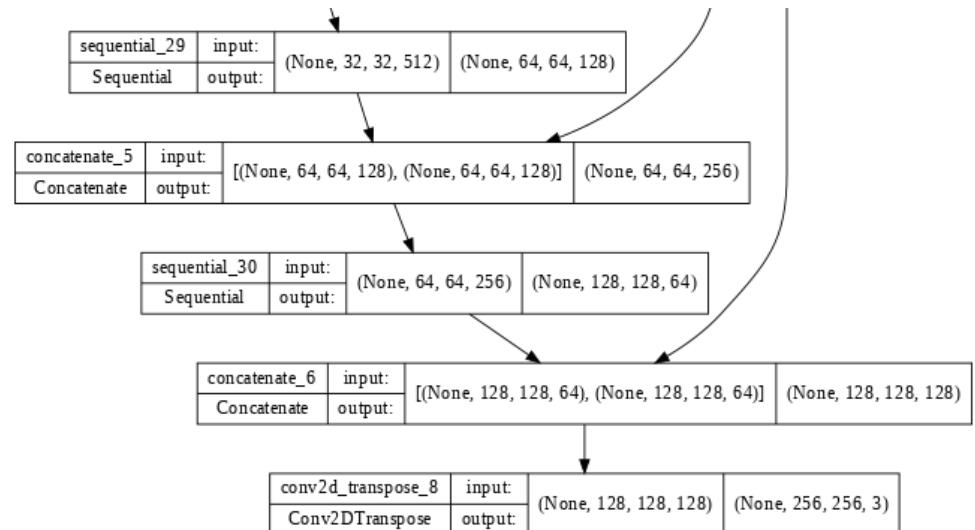
[Mostrar diferenças](#)



Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

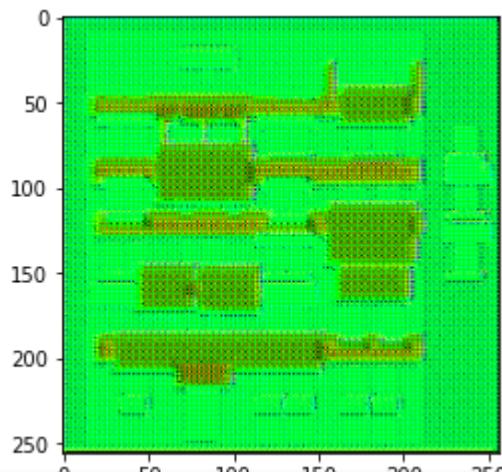
[Mostrar diferenças](#)





```
gen_output = generator(inp[tf.newaxis, ...], training=False)
plt.imshow(gen_output[0, ...])
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or  
<matplotlib.image.AxesImage at 0x7f85fd21fed0>



Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

LAMBDA = 100

```
loss_object = tf.keras.losses.BinaryCrossentropy(from_logits=True)
```

```
def generator_loss(disc_generated_output, gen_output, target):
```

```
gan_loss = loss_object(tf.ones_like(disc_generated_output), disc_generated_output)

# Mean absolute error
l1_loss = tf.reduce_mean(tf.abs(target - gen_output))

total_gen_loss = gan_loss + (LAMBDA * l1_loss)

return total_gen_loss, gan_loss, l1_loss

def Discriminator():
    initializer = tf.random_normal_initializer(0., 0.02)

    inp = tf.keras.layers.Input(shape=[256, 256, 3], name='input_image')
    tar = tf.keras.layers.Input(shape=[256, 256, 3], name='target_image')

    x = tf.keras.layers.concatenate([inp, tar])  # (batch_size, 256, 256, channels*2)

    down1 = downsample(64, 4, False)(x)  # (batch_size, 128, 128, 64)
    down2 = downsample(128, 4)(down1)  # (batch_size, 64, 64, 128)
    down3 = downsample(256, 4)(down2)  # (batch_size, 32, 32, 256)

    zero_pad1 = tf.keras.layers.ZeroPadding2D()(down3)  # (batch_size, 34, 34, 256)
    conv = tf.keras.layers.Conv2D(512, 4, strides=1,
                                kernel_initializer=initializer,
                                use_bias=False)(zero_pad1)  # (batch_size, 31, 31, 512)

    batchnorm1 = tf.keras.layers.BatchNormalization()(conv)

    leaky_relu = tf.keras.layers.LeakyReLU()(batchnorm1)

    zero_pad2 = tf.keras.layers.ZeroPadding2D()(leaky_relu)  # (batch_size, 33, 33, 512)

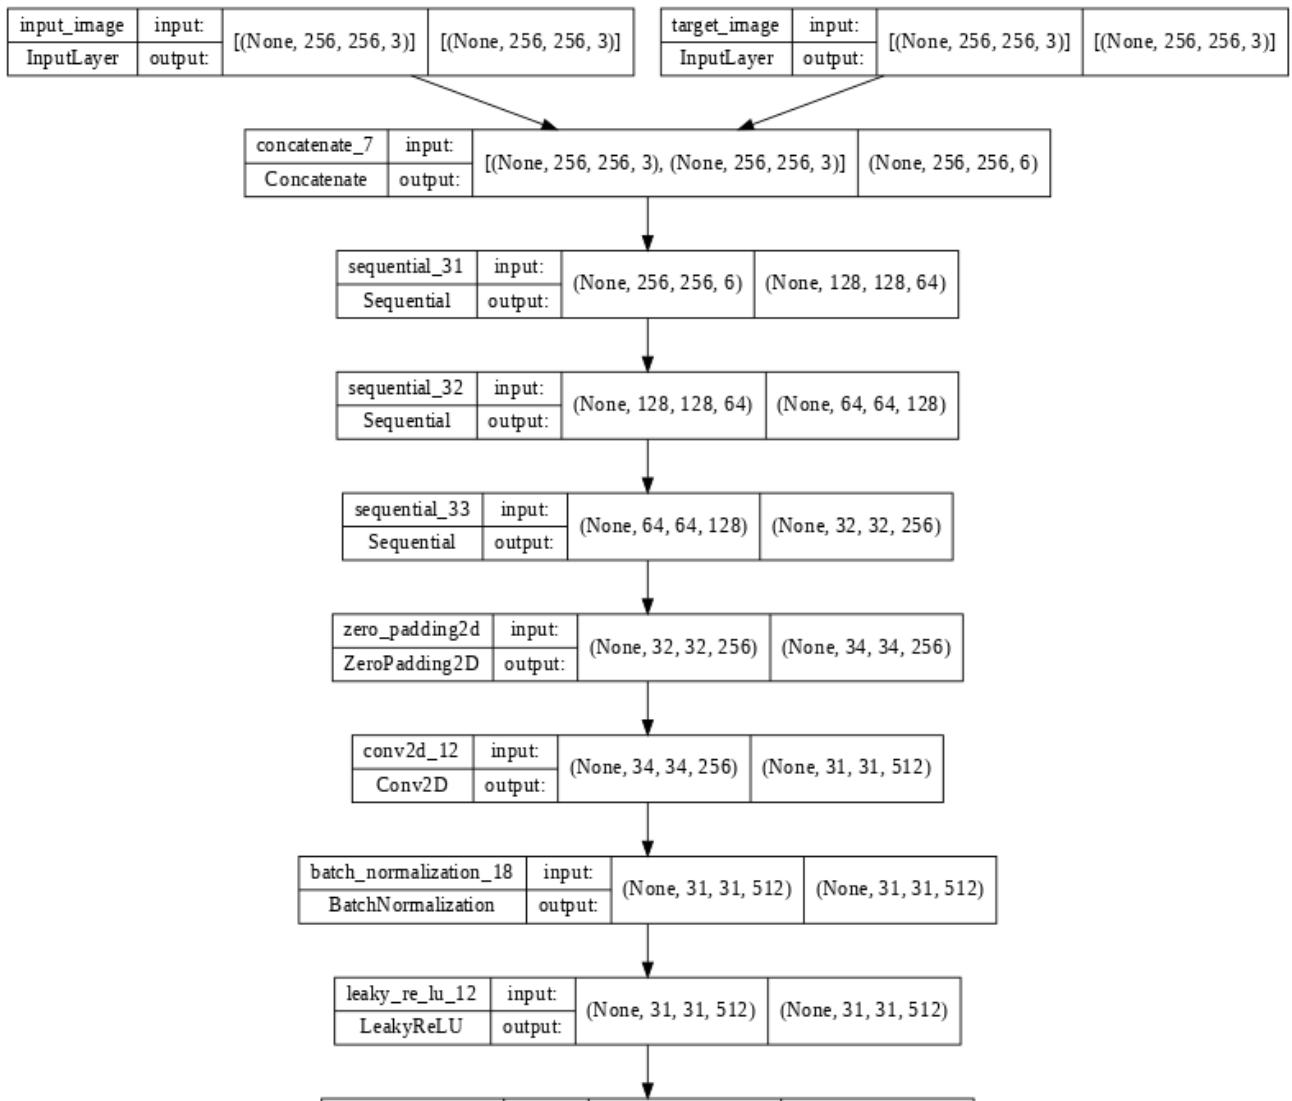
    last = tf.keras.layers.Conv2D(1, 4, strides=1,
                                kernel_initializer=initializer)(zero_pad2)  # (batch_size, 30, 30, 1)

    return tf.keras.Model(inputs=[inp, tar], outputs=last)

discriminator = Discriminator()
tf.keras.utils.plot_model(discriminator, show_shapes=True, dpi=64)
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

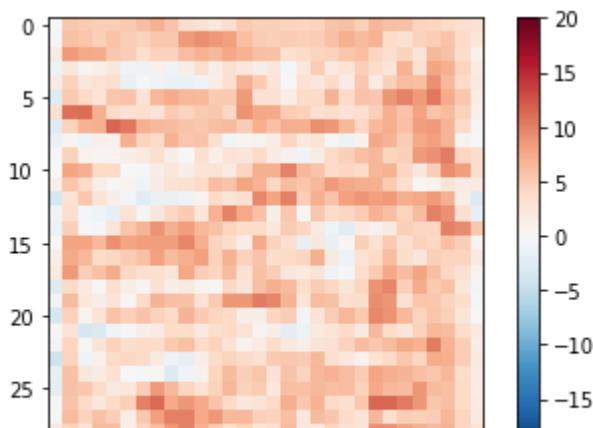
[Mostrar diferenças](#)



```

disc_out = discriminator([inp[tf.newaxis, ...], gen_output], training=False)
plt.imshow(disc_out[0, ..., -1], vmin=-20, vmax=20, cmap='RdBu_r')
plt.colorbar()
  
```

<matplotlib.colorbar.Colorbar at 0x7f86205b0f10>



Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```

def discriminator_loss(disc_real_output, disc_generated_output):
    real_loss = loss_object(tf.ones_like(disc_real_output), disc_real_output)

    generated_loss = loss_object(tf.zeros_like(disc_generated_output), disc_generated_output)
  
```

```

total_disc_loss = real_loss + generated_loss

return total_disc_loss

generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)

checkpoint_dir = './training_checkpoints'
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
                                 discriminator_optimizer=discriminator_optimizer,
                                 generator=generator,
                                 discriminator=discriminator)

def generate_images(model, test_input, tar):
    prediction = model(test_input, training=True)
    plt.figure(figsize=(15, 15))

    display_list = [test_input[0], tar[0], prediction[0]]
    title = ['Input Image', 'Ground Truth', 'Predicted Image']

    for i in range(3):
        plt.subplot(1, 3, i+1)
        plt.title(title[i])
        # Getting the pixel values in the [0, 1] range to plot.
        plt.imshow(display_list[i] * 0.5 + 0.5)
        plt.axis('off')
    plt.show()

for example_input, example_target in test_dataset.take(1):
    generate_images(generator, example_input, example_target)

```



Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
log_dir="logs/"
```

```
summary_writer = tf.summary.create_file_writer(
    log_dir + "fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
```

```

@tf.function
def train_step(input_image, target, step):
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        gen_output = generator(input_image, training=True)

        disc_real_output = discriminator([input_image, target], training=True)
        disc_generated_output = discriminator([input_image, gen_output], training=True)

        gen_total_loss, gen_gan_loss, gen_l1_loss = generator_loss(disc_generated_output, gen_disc_loss = discriminator_loss(disc_real_output, disc_generated_output)

        generator_gradients = gen_tape.gradient(gen_total_loss,
                                                generator.trainable_variables)
        discriminator_gradients = disc_tape.gradient(disc_loss,
                                                       discriminator.trainable_variables)

        generator_optimizer.apply_gradients(zip(generator_gradients,
                                                generator.trainable_variables))
        discriminator_optimizer.apply_gradients(zip(discriminator_gradients,
                                                    discriminator.trainable_variables))

    with summary_writer.as_default():
        tf.summary.scalar('gen_total_loss', gen_total_loss, step=step//1000)
        tf.summary.scalar('gen_gan_loss', gen_gan_loss, step=step//1000)
        tf.summary.scalar('gen_l1_loss', gen_l1_loss, step=step//1000)
        tf.summary.scalar('disc_loss', disc_loss, step=step//1000)

def fit(train_ds, test_ds, steps):
    example_input, example_target = next(iter(test_ds.take(1)))
    start = time.time()

    for step, (input_image, target) in train_ds.repeat().take(steps).enumerate():
        if (step) % 1000 == 0:
            display.clear_output(wait=True)

            if step != 0:
                print(f'Time taken for 1000 steps: {time.time()-start:.2f} sec\n')

            start = time.time()

            generate_images(generator, example_input, example_target)
            print(f"Step: {step//1000}k")

        train_step(input_image, target, step)

```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
print('. ', end='', flush=True)
```

```
# Save (checkpoint) the model every 5k steps
if (step + 1) % 5000 == 0:
    checkpoint.save(file_prefix=checkpoint_prefix)
```

```
%load_ext tensorboard  
%tensorboard --logdir {log_dir}
```

TensorBoard INACTIVE

---

**No dashboards are active for the current data set.**

Probable causes:

- You haven't written any data to your event files.
- TensorBoard can't find your event files.

If you're new to using TensorBoard, and want to find out how to add data and set up your event files, check out the [README](#) and perhaps the [TensorBoard tutorial](#).

If you think TensorBoard is configured properly, please see [the section of the README devoted to missing data problems](#) and consider filing an issue on GitHub.

*Last reload: Jun 1, 2022, 1:14:36 PM*

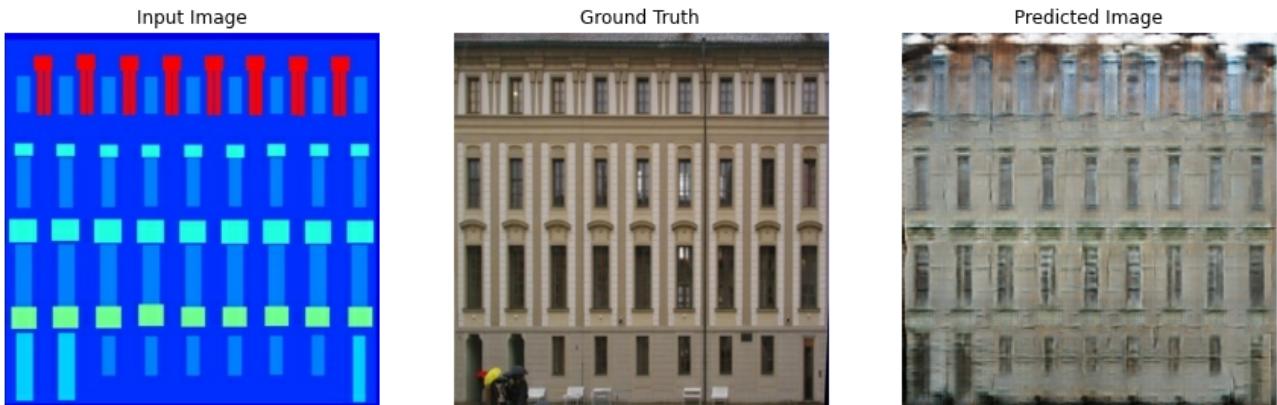
*Log directory: logs/*

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
fit(train_dataset, test_dataset, steps=40000)
```

Time taken for 1000 steps: 94.75 sec



Step: 39k

```
tensorboard dev upload --logdir {log_dir}
```

ERROR: Failed to launch TensorBoard (exited with 1).

Contents of stderr:

\*\*\*\*\* TensorBoard Uploader \*\*\*\*\*

This will upload your TensorBoard logs to <https://tensorboard.dev/> from the following directory:

logs/

This TensorBoard will be visible to everyone. Do not upload sensitive data.

Your use of this service is subject to Google's Terms of Service

<<https://policies.google.com/terms>> and Privacy Policy

<<https://policies.google.com/privacy>>, and TensorBoard.dev's Terms of Service

<<https://tensorboard.dev/policy/terms/>>.

This notice will not be shown again while you are logged into the uploader. To log out, run `tensorboard dev auth revoke`.

Traceback (most recent call last):

```
File "/usr/local/bin/tensorboard", line 8, in <module>
    sys.exit(run_main())
File "/usr/local/lib/python3.7/dist-packages/tensorboard/main.py", line 46, in run
    app.run(tensorboard.main, flags_parser=tensorboard.configure)
File "/usr/local/lib/python3.7/dist-packages/absl/app.py", line 312, in run
    _run_main(main, args)
File "/usr/local/lib/python3.7/dist-packages/absl/app.py", line 258, in _run_main
    sys.exit(main(argv))
File "/usr/local/lib/python3.7/dist-packages/tensorboard/program.py", line 276, in
    return runner(self.flags) or 0
File "/usr/local/lib/python3.7/dist-packages/tensorboard/uploader/uploader_subcomm
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

```
_print_to_user(Intent.FLAG_ACTIVITY_CLEAR_TOP)
File "/usr/local/lib/python3.7/dist-packages/tensorboard/uploader/uploader_subcomm
    response = input("Continue? (yes/NO) ")
EOFError: EOF when reading a line
Contents of stdout:
Continue? (yes/NO)
```

```
display.IFrame(  
    src="https://tensorboard.dev/experiment/lZ0C6FONROaUMfjYkVyJqw",  
    width="100%",  
    height="1000px")
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

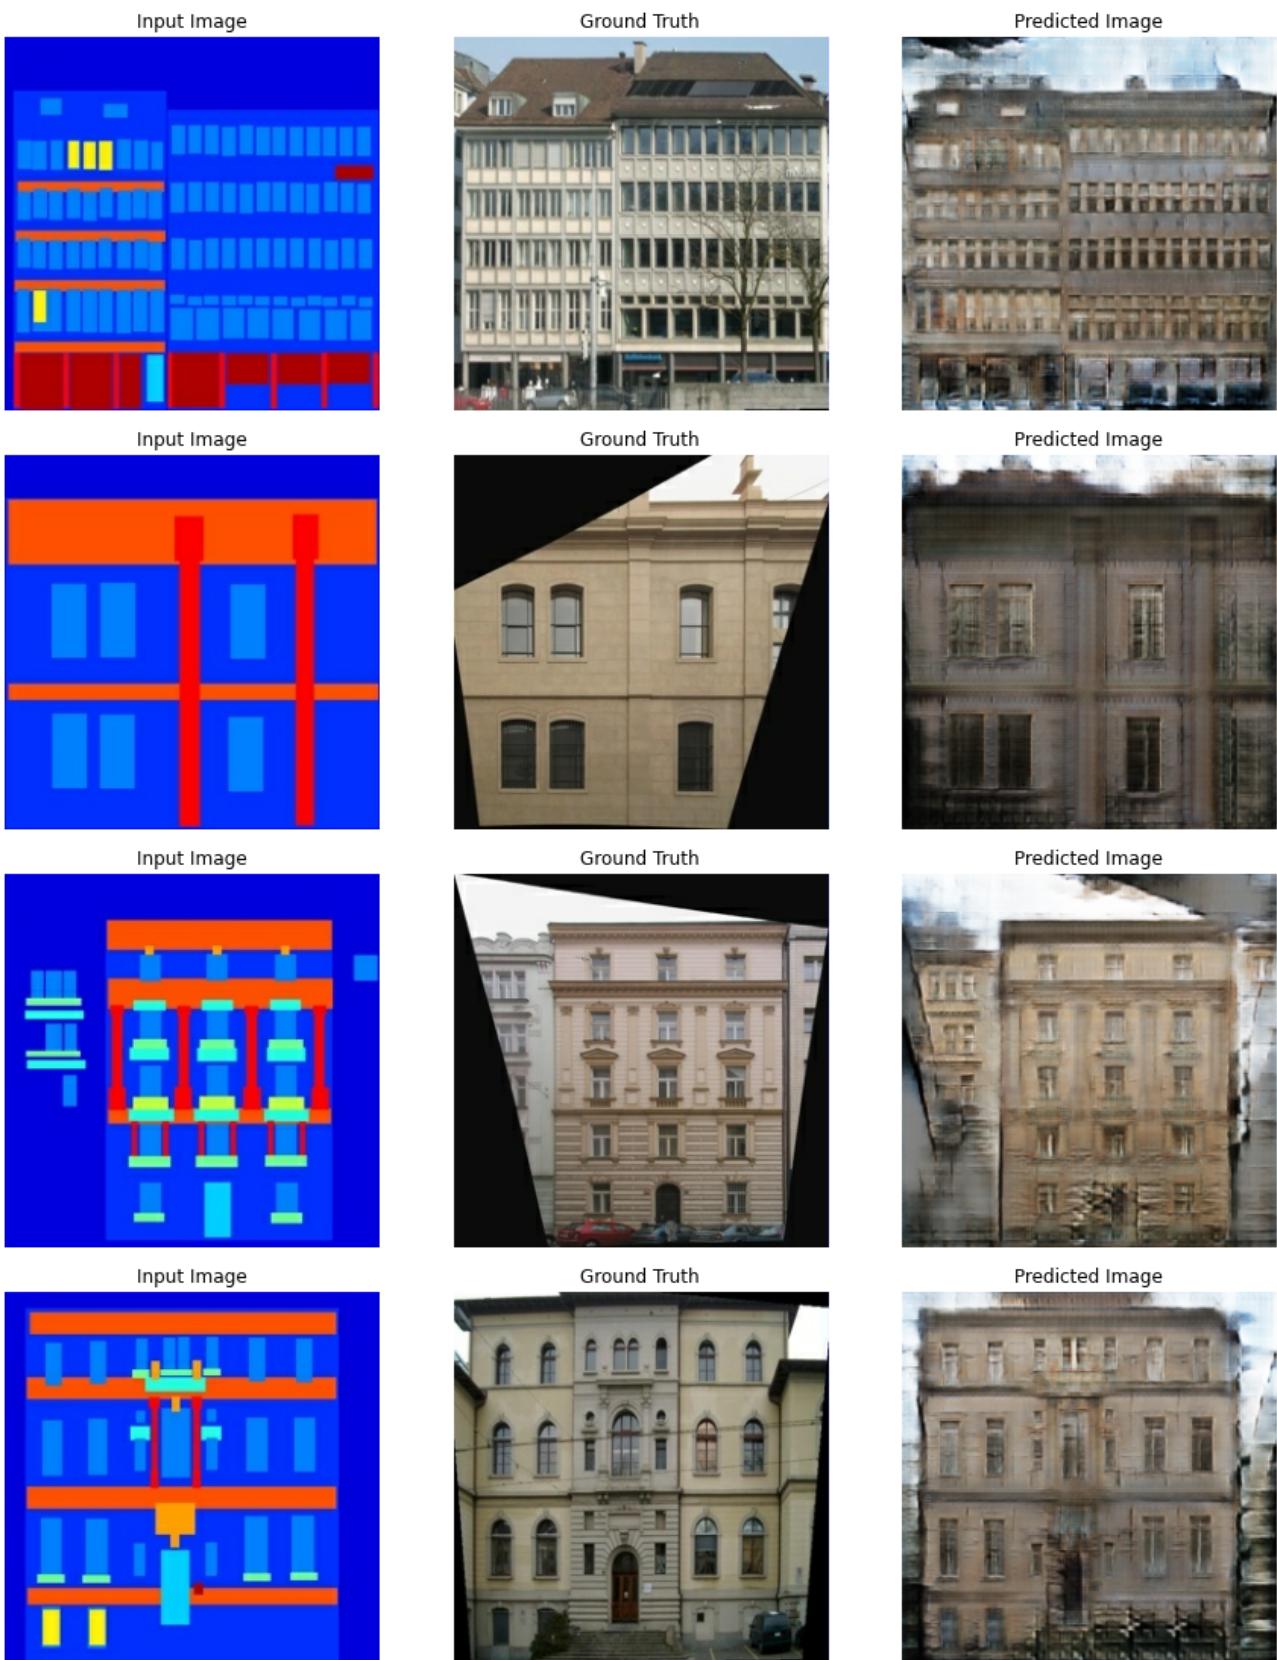
TensorBoard.dev >

[SEND FEEDBACK](#)

## SIGN IN

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra quia.

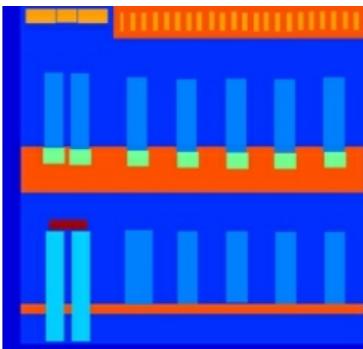
### Mostrar diferencias



```
for inp, tar in test_dataset.take(10):
```

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)



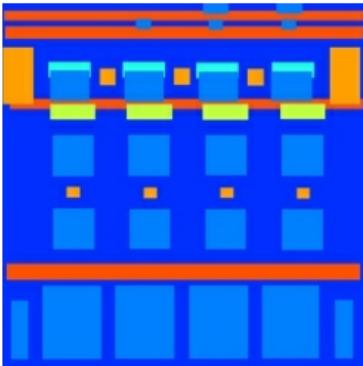
Input Image



Ground Truth



Predicted Image



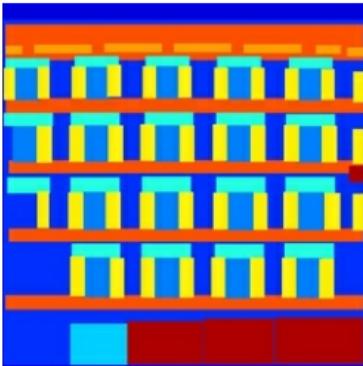
Input Image



Ground Truth



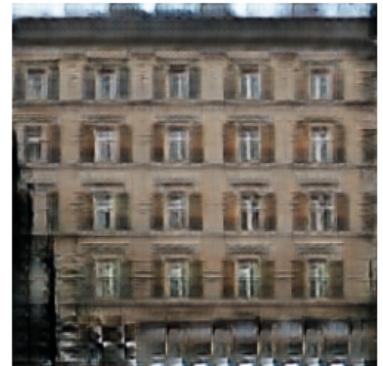
Predicted Image



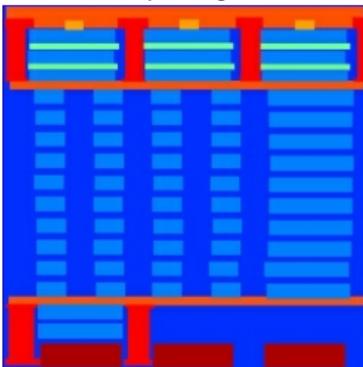
Input Image



Ground Truth



Predicted Image



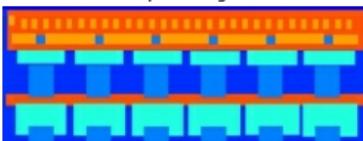
Input Image



Ground Truth



Predicted Image



Input Image



Ground Truth



Predicted Image

Falha no salvamento automático. Este arquivo foi atualizado remotamente ou em outra guia.

[Mostrar diferenças](#)

